

1 Problem Statement

Write a parallel program in C/C++ to estimate the value of π using the Monte Carlo method specified below. The program should take n as input, which is the number of points to be used for the estimation. The program should then generate n random points in the unit square $[0, 1] \times [0, 1]$ and count the number of points that lie inside the unit circle.

Let n be a large integer, then estimated value of π is

$$\pi \approx \frac{4 \cdot \text{number of points in the circle}}{n}$$




Figure 1: Estimating π using Monte Carlo method

2 Parallel Algorithm

The estimation can be easily parallelized by assigning each processor the responsibility for $\frac{n}{p}$ distinct points. You should use the following MPI functions:

- Have processor with rank 0 read n , and use `MPI_Bcast` to broadcast it to all processors.
- Use `MPI_Reduce` function to sum the number of points in the circle across all processors.

- Use `MPI_Wtime` to time the run-time of the program (measure on processor 0).

For the random number generator, you can use the `rand()` function in C/C++. However, you should make sure that the random number generator is seeded **differently** for each processor. You can use the processor rank for this purpose. For example, if the rank of the processor is `i`, then you can use `srand(time(NULL) + i)` to seed the random number generator on that processor.

3 Code framework

3.1 Input & Output Format

Your program should take n as the input using command line arguments and output the estimated value of π and the time taken to compute this value (comma-separated). For example, if the estimated value is `s` and the time taken is `t`, your output should be “`s, t`”. (**Note:** `s` should be printed up to 12 decimal points). All output is done through processor with rank 0.

3.2 Deliverables

1. Create a Makefile for your program, and make sure the name of your output executable is “`pi_calc`”. If you are not familiar with creating Makefiles, check resources below for help.
2. Write a “`README.txt`” briefly describing how your program works and the machine you used for generating the results.
3. For $n = 10^6$, plot a graph of run-time of the program vs. the number of processors for a few chosen values of p . This run-time should include all computations that contribute to the estimation, including any local computations and global reductions. Include your graph and observation regarding the speedup in a PDF file with name “`report.pdf`”. **Make sure to list names of all your teammates at the very beginning of your report.**
4. Submit your a) “`code.zip`” in “Programming Assignment 1 - **Code**” on Gradescope. Your zip file should include all the `.cpp` files, `Makefile` and `README.txt` file; b) your report in “Programming Assignment 1 - **Report**” on Gradescope.

4 Resources

1. What is a Makefile and how does it work?: <https://opensource.com/article/18/8/what-how-makefile>
2. PACE ICE cluster guide: https://docs.pace.gatech.edu/ice_cluster/ice-guide/. Documentation for writing a PBS script: https://docs.pace.gatech.edu/software/PBS_script_guide/